

RISC-V Processor Design Literature Review and Research Proposal Formulation

Matt Young
s4697249
m.young2@uqconnect.edu.au

November 2023

Contents

1	Introduction	1
2	Methodology	2
3	Initial ideas	3
3.1	A triple-redundancy fault tolerant RISC-V processor for safety-critical applications	3
3.2	A VLIW RISC-V processor for embedded DSP applications	4
3.3	A runtime reconfigurable microarchitecture for RISC-V processors	5
4	Literature reviews	6
4.1	Reconfigurable computing	6
4.1.1	Data compression	6
4.1.2	Course Grained Reconfigurable Architectures (CGRAs)	7
4.1.3	Dataflow architectures	7
4.1.4	Approximate computing	7
4.1.5	SLAM	8
5	Industry trends	9
6	Finalised research proposals	9
7	References	9

1 Introduction

This document is a follow on from “Project Proposal: Design, Verification and Synthesis of a RISC-V RV32IC Processor”, which we discussed in early October 2023. As we discussed at the time, that project was interesting in a practical sense, but it needed to address a specific *research question* to be viable as a thesis. This document aims to find that research question, both by detailing ideas that I myself have had over the past few years, as well as performing a relatively detailed literature review of recent papers published in the fields of RISC-V processor design and reconfigurable computing.

Before we set out to find a specific research question, we should ask ourselves: What is an acceptable thesis project? My aim is to make this project the best it can possibly be, but the timeline of one year and my relative lack of experience in processor design introduces some constraints. Specifically, we want the project to be complex enough to explore relatively novel or non-traditional CPU design ideas, but we also need to make sure the project can be completed and written up comfortably in one year. Unfortunately, this particular constraint means I have to rule out some interesting ideas such as anything out-of-order, speculative or superscalar. It does also raise into question whether we attempt more complicated things such as multi/many-core designs (even if backed by a relatively simple per-core architecture), and VLIW. Nonetheless, these complicated topics are something I'm looking forward to considering in future research, but for now, I think we are best finding a concept that is simple in theory and can be safely executed in one year (considering other coursework as well). If the project does end up being too simple, we can always add extensions.

2 Methodology

To perform the literature review, we will consider both academia as well as industry. The aim is to determine what the latest research in academia is, and to compare that with what the latest industry trends are.

The specific topics we will look into are:

- RISC-V CPU architectures, preferably simpler ones
- Low-power processor design
- Very-long-instruction-word
- Branch prediction algorithms, preferably simple, low-power ones
- Reconfigurable computing
- Custom instructions vs. Custom hardware, for domain-specific acceleration
- Applications of custom processors and/or RISC-V in cybersecurity

These topics will primarily be sourced from the latest editions of the following journals:

- IEEE International Conference on Field-Programmable Technology (FPT)
- Annual IEEE Symposium on Field-Programmable Custom Computing Machines (FCCM)
- Reconfigurable Architectures Workshop

Finally, note that we will exclude anything regarding AI and machine learning, purely because this is not a topic I'm interested in researching. That being said, a lot of AI-centric designs can be repurposed to do more interesting DSP work, so I won't completely discard AI papers if they seem to be useful for DSP.

Before we start the literature review, however, I will first briefly detail some ideas I personally have come up with over the last two years or so, which might serve as useful research questions. The goal of the literature review, then, is to see which of these ideas are reasonable, as well as to discover new topics which I might not have considered yet. Once this is done, we will end up with a list of research questions which meet all the relevant criteria for my thesis!

3 Initial ideas

The following section details various paper ideas I’ve had over the last few years. The goal is to lay these out first, and then evaluate if they hold up to scrutiny as actually good ideas by performing the literature review. The ideas appear in the order of how much I think they are a reasonable project, with the most reasonable first and least reasonable last. Each idea has a suggested thesis title and abstract, along with pros/cons, and some additional notes, and finally a concluding verdict.

Note that these are only the ideas I, personally, have come up with. In the process of performing the literature review, we will undoubtedly come across many more good ideas that I couldn’t ever think of. We will catalogue both the ideas I’ve had, and ideas inspired by the literature review, in a final section.

3.1 A triple-redundancy fault tolerant RISC-V processor for safety-critical applications

Abstract: In recent years, there has been increasing interest in designing RISC-V processors for safety-critical applications. For example, NASA selected SiFive to provide a RISC-V processor for their High-Performance Spaceflight Computing project. One way to design a fault tolerant processor is by using one or more redundant cores operating in lockstep. In this thesis, I design a RISC-V RV32ICM processor with three cores in a triply-redundant lockstep configuration. The design includes attractive features for embedded computing such as a single-cycle 32x32 multiplier and a barrel shifter. The design is synthesised for the Lattice ECP5 FPGA using open-source tools, and the resulting physical implementation is benchmarked using CoreMark. Combined with a rad-hardened CMOS fabrication node, this processor may be an attractive offering for mission-critical computing in defence or space applications, or as an FPGA soft-core for industrial applications.

Pros:

- Interesting research topic, as I am interested in multi-/many-core designs
- Definite real-world applications, probably more so as a softcore rather than an ASIC
- The same processor can effectively be duplicated 3 times, it’s just the redundancy and potentially correction logic that is hard

Cons:

- Redundancy logic may be difficult to verify, difficult to prove it’s actually redundant (I imagine a lot of edge cases)
- We will probably encounter significant difficulties with verification, especially with fault-tolerance.
- We may encounter significant time-pressures as designing a RV32 core is complex enough, and the fault tolerant logic adds extra verification work.

Notes:

- We need to *very* carefully scope out this project and define *exactly* what we mean by “fault tolerant”. This will most likely involve designing to be immune to certain types of single event upsets.

- We need to research ways we can reliably introduce single event upsets into both our simulated and real-world system to prove fault-tolerance. This may be extremely challenging.
- A quick search reveals that there are many existing papers on fault tolerant RISC-V design - we could learn a lot from these.

Verdict: I think this is a good project, it's achievable and has backup options. We should explore research and strongly consider this idea. Specifically, we need to research the why and how of fault tolerant hardware design, see if anyone's made a fault-tolerant RISC-V processor (and if so, how they did it), and also particularly understand triple-redundancy.

3.2 A VLIW RISC-V processor for embedded DSP applications

Abstract: Digital signal processing (DSP) has become an increasingly important workload for edge computing devices, such as microcontrollers and low-power microprocessors. Yet, maintaining high DSP performance while managing power remains a difficult challenge. VLIW processors have a long history in the DSP segment of the semiconductor industry for their ability to achieve high performance while requiring only simple, low-power hardware. In this thesis, I explore designing a very-long-instruction-word (VLIW) RISC-V processor with specific extensions for fixed-point calculations, ideal for DSP workloads. The processor uses a special VLIW encoding based on the RISC-V RV32IM ISA, and processes 128-bit words for a maximum of 4 issues per clock. Each of the 4 execution unit supports a single-cycle 32x32 multiplier. I synthesise this design for the Lattice ECP5 FPGA using open-source tools, and measure its power, performance and area (PPA) against competing RISC-V processors on a simple DSP workload.

Pros:

- Very interesting research: combines bits of my two favourite interests, processor design and high-performance computing
- Some potential real-world applications, although not as much as other topics
- Relatively novel research, there are not many papers on RISC-V VLIW encodings
- Strictly speaking, VLIW hardware is not “that” hard (it's just the writing the decoder and verifying it that's the hard part - the functional units can be easily duplicated)

Cons:

- Harder to verify: unlike a traditional RV32IC core, the encoding is semi-custom, so we can't use any existing RISC-V verification suites or instruction fuzzers
- May encounter significant difficulties with the register file, especially since the ECP5 only supports a maximum of dual-port RAM, where we may need quad-port or higher for VLIW
- No RISC-V VLIW compilers (and writing a compiler out of scope for this thesis), so we would have to write an assembler and then program DSP algorithms in assembly - ouch!
- Although VLIW is cool (in my opinion), it's not as applicable in industry as a more standard RV32 design

Notes:

- We need to clarify exactly what is meant by “fixed point extensions” - what hardware is involved in this? Do we want a CORDICS module?

Verdict: Explore more research. This may be too hard for us right now.

3.3 A runtime reconfigurable microarchitecture for RISC-V processors

Abstract: Modern processors are designed with power, performance and area (PPA) in mind. However, at times, these goals can be conflicting: a low-power design naturally reduces performance, and a high-performance design increases power. In this thesis, I introduce a RISC-V RV32IC processor with a runtime reconfigurable microarchitecture. End-users can configure the processor’s microarchitecture at runtime to dynamically target the goals of high-performance or low-power, with a great deal of flexibility. This is achieved by installing both a “high-performance” and “low-power” microarchitecture on the same chip, and using clock-gating and a re-routing system based on writes to a specific register. The design is synthesised for the Lattice ECP5 FPGA using open-source tools, and real-world power and performance metrics are analysed.

Pros:

- Interesting research topic
- Potential real-world applications and future research options

Cons:

- Unclear if my ideas for the low-power side of things will actually work
- In particular, the multi-cycle design from Harris & Harris may actually draw more power than a pipelined design. This would need some serious analysis.
- Requires essentially implementing two processors: high performance and low power

Notes:

- Instead of implementing two CPUs from scratch, we could borrow existing implementations - this is less fun though
- Following Harris & Harris might naturally allow us to come up with the “low power” and “high performance” designs
- How does this compare to big.LITTLE style designs (i.e. heterogeneous computing)? I discarded the idea of making a big.LITTLE style CPU because, in my opinion, it doesn’t make sense to do for a microcontroller. Does a “reconfigurable microarchitecture” make any sense on a microcontroller either? We may not see much power difference unless you compare something like an OoO design with a in-order design, and of course OoO is *way* out of our league.

Verdict: This is an interesting question, but it may be hard both to implement, and we may not see any real power savings. It also may not make sense to do for a microcontroller style CPU anyway.

4 Literature reviews

This brief literature review aims to determine what the latest research in both reconfigurable computing and processor design is. The reconfigurable computing journals I will be analysing are the IEEE FPT and FCCM, as well as the RAW workshop (not from IEEE). On the processor design side, the journals that will be analysed are:

4.1 Reconfigurable computing

As might be expected, accelerating compute-intensive workloads remains one of the principal applications of reconfigurable computing. However, the most “trendy” workload to accelerate at the moment appears to be machine learning, specifically deep learning. In my opinion, I believe that there are more devices actually deployed in the real-world that need DSP accelerating,¹ however, undoubtedly the more popular research topic is machine learning. Although DSP remains an incredibly important workload for FPGAs, machine learning is being added to more and more devices, so will become an increasingly important workload as years go on. Already, we’ve seen manufacturers like Xilinx offer AI-tuned SoCs and IP cores.

4.1.1 Data compression

Other workloads that remain popular for FPGA acceleration include data compression, as shown in [1]. The authors in this particular paper demonstrate performance improvements of up to 25x a fast AMD CPU on a challenging compression algorithm. FPGAs are a great target for data compression due to their hardware parallelism. We could potentially investigate adding data compression instructions to a RISC-V processor, in particular regarding zlib, bzip2 or zstd. The algorithm tested in [1] seems to be too niche, and is complicated to accelerate compared to something like zlib or DEFLATE. We should also consider comparing whether an FPGA-SoC like Zynq, or processor with custom compression functions included, is faster. There is some existing literature on this concept: In [2], researchers from IBM demonstrate their hardware data compression accelerator for the taped-out IBM POWER9 processor. This circuit is designed into the main silicon of the chip, and the entire processor itself is manufactured on GlobalFoundries’ 14nm process node. Their results indicate that the compressor circuit uses only 0.5% chip area, but can speed up zlib workloads 388x compared to a single IBM POWER9 core. This is extremely impressive, as POWER9 is a modern, speculative out-of-order processor [3].

In terms of applying this concept to RISC-V processors, there is exceedingly little research. The only recent paper which does so was published by Google researchers in [4], in which, as part of their benchmarking suite, they introduce the concept of a “compression and decompression processing unit”, or “CDPU”. They develop a “CDPU” RTL generator in Chisel [5], and extend the famous out-of-order BOOM core [6] from Berkeley with their new CDPU. They report that their CDPU accelerator is 10x to 16x faster than a single Xeon core at various compression algorithms, including some new ones like zstd. If we were to implement a data compression accelerator, we would most certainly not be able to target anything this complicated. We could instead focus on simpler algorithms, such as Huffman coding or even just run-length encoding, to demonstrate the principle that compression can be accelerated in hardware. This topic would yield pretty good results, however, it unfortunately misses out on the processor *design*, as we would most likely be using someone else’s RISC-V core and just adding a compression accelerator to it.

¹I have no actual evidence to back this up other than a gut feeling.

4.1.2 Course Grained Reconfigurable Architectures (CGRAs)

One new area of research in reconfigurable computing is called a “course grained reconfigurable architecture”, or CGRA. The goal of a CGRA, compared to an FPGA, is to significantly improve performance at the cost of the granularity of reconfiguration. Compared to today’s FPGAs, which almost exclusively use SRAM-based lookup tables (LUTs) and a reconfigurable fabric, CGRAs are a newer and less standardised technology, so there is not one simple general CGRA architecture that can be pointed to. In [7], the authors perform an extremely detailed literature review of existing CGRA designs. They state that modern CGRAs typically consist of a number of functional units, each which contain an ALU, possibly a multiplier, and a register file. These FUs are then networked together in a reconfigurable mesh, similar to an FPGA. When techmapping an algorithm onto a CGRA processor, typically the innermost loop of the algorithm will be selected. This works extremely well for DSP-style applications, such as IIR filters or FFTs. (TODO cite that VLIW paper/book) However, as stated in [8] and [9], this process is extremely complex and time consuming for the compiler, as mapping the algorithm’s Data Dependency Graph onto the target CGRA is NP-complete.

CGRAs specifically came to my attention because of the article [10] published in the FPT journal. In this paper, the authors propose several CGRA architectures to reduce connections across what they call switch blocks. They also analyse the hardware cost and routability of various data-flow graphs, to see how algorithms would map onto the proposed CGRA.

While CGRAs are extremely interesting, they may not be the best topic of research for this particular thesis. Designing a CGRA would be a very interesting research topic, however, it would require a VLSI workflow and even possibly a tapeout to see genuine performance gains. Like other digital circuits, it’s entirely possible to prototype it on an FPGA, however, the performance degradation will be significant. Without a VLSI flow, which is not currently available at UQ, we would only be able to determine expected performance improvements based on RTL simulation. Since CGRAs are relatively cutting-edge, acquiring the necessary skills to develop a novel CGRA implementation would also be difficult. With all that said, CGRAs are certainly an interesting topic for future research.

4.1.3 Dataflow architectures

4.1.4 Approximate computing

Recently, the concept of approximate computing has emerged and become an actively researched topic in computer engineering. One upcoming industry trend is the deployment of neural network machine learning models to edge computing devices, such as smart cameras or home assistants. The perceptron model that these neural networks are based on does not need particularly accurate computing requirements in order to function with a high degree of similarity to full-precision models. Currently, most machine learning frameworks such as Tensorflow and PyTorch offer the ability to heavily quantise the neuron weights, for example to an 8-bit fixed point format. This means that many pieces of machine learning hardware such as Google’s Tensor Processing Units (TPUs) are optimised for large scale, low precision arithmetic. However, with approximate computing, this quantisation can be made more efficient. Approximate adders and multipliers, which use less power and area than their full-precision counterparts, can be combined with quantisation to produce extremely efficient edge computing hardware capable of executing the latest machine learning models quickly while drawing minimal power.

The RISC-V ISA is an excellent platform to extend with approximate computing extensions. There already exist a number of RISC-V machine learning accelerators, and these could be further improved by designing and implementing an approximate computing extension. This

may have applications not just in machine learning, but also possibly in fields where high-volume, low-precision compute is required, such as certain types of embedded DSP. We could design a processor that has a particular register, or perhaps a custom instruction, that switches from a precise ALU to an approximate ALU. Alternatively, we could add an additional set of instructions for the RISC-V ISA that uses approximate computing, so that precise and approximate arithmetic can be mixed.

4.1.5 SLAM

Simultaneous localisation and mapping, or SLAM, is a very important technique in robotics for both determining the robots current position (“localisation”) and producing either a 2D or 3D model of the environment (“mapping”). SLAM algorithms can be run on vision, LiDAR, radar and other sensors. Because of the pattern matching and linear algebra involved in SLAM algorithms, they can be rather expensive to run. Currently, as of writing, I’m employed as an intern at Emesent, a robotics company which produces Hovermap, a leading drone-based GPS-denied 3D LiDAR scanning platform. We use an extremely powerful SLAM algorithm called Wildcat [11], originally developed at CSIRO. Although Wildcat makes use of efficient linear algebra primitives, running the algorithm on the embedded compute of the Hovermap payload can still present a challenge. I have often thought that using a Zynq SBC and offloading some of the SLAM work to an FPGA might be able to improve performance. In [12], the authors develop “FSLAM”, an FPGA accelerated implementation of the ORB-SLAM visual SLAM algorithm. They report that their implementation is 8.5x faster than an ARM CPU, and 1.55x faster than an “Intel desktop CPU”². In table IV in the paper, their FPGA implementation achieves 62 FPS and draws only 4.6W of power. Although the desktop GPU implementation is the fastest, at 70 FPS, it draws 200W of power (!)

Overall, this SLAM accelerator idea is a very interesting research topic. I have thought about whether processor designs optimised for DSP, such as VLIW processors, would be applicable for SLAM acceleration. Realistically, although more performance might be gained out of a specialised processor, for LiDAR SLAM acceleration, you would want to use an FPGA, as porting Wildcat plus all the libraries it uses to RISC-V would be extremely challenging. In addition, the existing implementation *likely* uses x86 SIMD instructions, which almost certainly yield more performance than the RISC-V vector extension. The better way to extract more performance would be adding specialised hardware units for direct point cloud feature extraction and using DMA. Even on an FPGA SoC like a Zynq, with extensive use of DMA and even PCI-Express interconnect, there is still no guarantee that memory wouldn’t be the bottleneck and we would receive worse performance than keeping it within the CPU. We would probably need some extensive caching mechanism to achieve decent performance.

I work with the lead developer of Wildcat, so I can ask him the possibility of accelerating it using an FPGA and how realistic of a proposal that is. If this does end up being a good idea, using my employment at Emesent, we could probably do this project in collaboration with them.

²Somewhat annoyingly, the authors do not state what specific CPU this is except for it being an Intel i5. We can presume it’s a modern OoO x86 core.

5 Industry trends

6 Finalised research proposals

7 References

- [1] Dongdong Tang et al. “pLPAQ: Accelerating LPAQ Compression on FPGA”. In: *2022 International Conference on Field-Programmable Technology (ICFPT)*. 2022, pp. 1–6. DOI: [10.1109/ICFPT56656.2022.9974593](https://doi.org/10.1109/ICFPT56656.2022.9974593).
- [2] Bulent Abali et al. “Data Compression Accelerator on IBM POWER9 and z15 Processors : Industrial Product”. In: *2020 ACM/IEEE 47th Annual International Symposium on Computer Architecture (ISCA)*. 2020, pp. 1–14. DOI: [10.1109/ISCA45697.2020.00012](https://doi.org/10.1109/ISCA45697.2020.00012).
- [3] Satish Kumar Sadasivam et al. “IBM Power9 Processor Architecture”. In: *IEEE Micro* 37 (2017), pp. 40–51. URL: <https://api.semanticscholar.org/CorpusID:31208323>.
- [4] Sagar Karandikar et al. “CDPU: Co-Designing Compression and Decompression Processing Units for Hyperscale Systems”. In: *Proceedings of the 50th Annual International Symposium on Computer Architecture*. ISCA '23. Orlando, FL, USA: Association for Computing Machinery, 2023. ISBN: 9798400700958. DOI: [10.1145/3579371.3589074](https://doi.org/10.1145/3579371.3589074). URL: <https://doi.org/10.1145/3579371.3589074>.
- [5] Jonathan Bachrach et al. “Chisel: Constructing Hardware in a Scala Embedded Language”. In: *Proceedings of the 49th Annual Design Automation Conference*. DAC '12. San Francisco, California: Association for Computing Machinery, 2012, pp. 1216–1225. ISBN: 9781450311991. DOI: [10.1145/2228360.2228584](https://doi.org/10.1145/2228360.2228584). URL: <https://doi.org/10.1145/2228360.2228584>.
- [6] Krste Asanović, David A. Patterson, and Christopher Celio. “The Berkeley Out-of-Order Machine (BOOM): An Industry-Competitive, Synthesizable, Parameterized RISC-V Processor”. In: 2015. URL: <https://api.semanticscholar.org/CorpusID:16283995>.
- [7] Artur Podobas, Kentaro Sano, and Satoshi Matsuoka. “A Survey on Coarse-Grained Reconfigurable Architectures From a Performance Perspective”. In: *IEEE Access* 8 (2020), pp. 146719–146743. DOI: [10.1109/ACCESS.2020.3012084](https://doi.org/10.1109/ACCESS.2020.3012084).
- [8] Shail Dave and Aviral Shrivastava. *Coarse-Grained Reconfigurable Arrays*. <https://labs.engineering.asu.edu/mps-lab/research/cgra/>. [Accessed 01-12-2023].
- [9] Compilers Creating Custom Processors (CCCP) Research Group. *CCCP: Coarse-Grained Reconfigurable Architecture*. <https://cccp.eecs.umich.edu/research/cgra.php>. [Accessed 01-12-2023]. 2016.
- [10] Boma Adhi et al. “Exploring Inter-tile Connectivity for HPC-oriented CGRA with Lower Resource Usage”. In: *2022 International Conference on Field-Programmable Technology (ICFPT)*. 2022, pp. 1–4. DOI: [10.1109/ICFPT56656.2022.9974525](https://doi.org/10.1109/ICFPT56656.2022.9974525).
- [11] Milad Ramezani et al. *Wildcat: Online Continuous-Time 3D Lidar-Inertial SLAM*. 2022. arXiv: [2205.12595](https://arxiv.org/abs/2205.12595) [cs.R0].
- [12] Vibhakar Vemulapati and Deming Chen. “FSLAM: an Efficient and Accurate SLAM Accelerator on SoC FPGAs”. In: *2022 International Conference on Field-Programmable Technology (ICFPT)*. 2022, pp. 1–9. DOI: [10.1109/ICFPT56656.2022.9974562](https://doi.org/10.1109/ICFPT56656.2022.9974562).