

Project Proposal: “Design, Verification and Synthesis of a RISC-V RV32IC Processor”

Matt Young
s4697249
m.young2@uqconnect.edu.au

May 2023

Contents

1	Executive summary (tl;dr)	1
2	Introduction	2
2.1	Background	2
2.2	This project	2
3	Microarchitecture design	2
3.1	ISA	2
3.2	Pipeline	3
3.3	Memory and privilege	3
3.4	Interrupts	3
3.5	Boot process	4
3.6	Other microarchitecture details	4
4	Tools and resources	4
4.1	HDL	4
4.2	Simulation, verification and synthesis	4
5	Verification plan	5
5.1	Verilator unit tests	5
5.2	RISC-V compliance test	5
5.3	Random instruction generation	5
6	Benchmarking	5
7	Possible extension tasks	5
7.1	Implement “speculative pre-fetching”	6
7.2	Add a fast “CLZ” instruction	6
7.3	Add “M” and “D” extensions	6
7.4	Implement a simple RISC pipeline	6
7.5	Add a simple branch predictor	6
7.6	SkyWater 130 nm VLSI tapeout using OpenLane	7

1 Executive summary (tl;dr)

1. I want to design, verify, and implement on an FPGA, a simple 32-bit RISC-V CPU core called “Jelly”. The CPU will be entirely designed and built by myself, from scratch.
2. The design’s goal is to be as simple as possible, while still allowing flexibility for future improvements.
3. Design, verification and logic synthesis will all be done using free and open-source tools.
4. I also already own the Lattice ECP5-85K FPGA that will be used, so I don’t need any funding or budget for this project - other than possibly access to electronics test equipment like a logic analyser.

5. This may possibly be the first CPU designed at UQ (or at least, in recent memory - as I was unable to find any prior art).
6. I believe I can achieve the above goals as I have experience with SystemVerilog and FPGAs already, I own all the tools required, and have been researching on this topic since about 2021.
7. There are some extension tasks such as adding a pipeline and multiply instructions that I may attempt if time permits.
8. Your involvement as a supervisor would be much appreciated!

2 Introduction

This document contains the initial proposal for my thesis project, “Design, Verification and Synthesis of a RISC-V RV32IC Processor”. Please note that this is only meant to explain the aim and goals of the project, not to be an official UQ thesis proposal.

2.1 Background

RISC-V is a popular, open-source, royalty-free, extensible instruction set architecture (ISA). It was originally designed by the University of California, Berkeley in 2015, but has since been managed by RISC-V International. Unlike ISAs such as ARM or x86, RISC-V is a fully open standard, meaning it can be implemented by a lowly computer science undergraduate without the risk of being sued.

Furthermore, RISC-V enjoys support from popular compiler toolchains, software vendors and the wider semiconductor industry. Many large companies such as NVIDIA, Western Digital, Seagate, Alibaba and Intel have taped out RISC-V processors that are currently in shipping products. There are also numerous RISC-V startups like SiFive, Tenstorrent, Esperanto and more who design new and novel processors.

Existing RISC-V processors can happily boot Linux, compile code with Clang and GCC, and even run complex applications like Firefox. Overall, RISC-V's wide industry adoption and open specification makes it a very attractive ISA to design a research processor around.

2.2 This project

The goal of my thesis is to design, verify and synthesise a simple 32-bit RISC-V CPU core (specifically one which implements the RV32IC ISA - which will be covered shortly). The implementation platform will be the Lattice ECP5-85K FPGA, which I already own via the OrangeCrab devkit.

This project has the codename “Jelly”, in reference to the jellyfish, which is apparently one of the simplest and most efficient animals, which are similar goals to this CPU.

3 Microarchitecture design

Jelly's goal is to provide as simple of a microarchitecture as reasonably possible, so that I can be certain I can get the project done in time, while still allowing flexibility for future upgrades and improvements.

The general design of the microarchitecture will be inspired by, although not directly lifted from, “Digital Design and Computer Architecture, RISC-V Edition” by Harris & Harris. This is essentially the seminal book on introductory RISC-V microarchitecture design, and will have many good references for my own design.

3.1 ISA

The specific architecture that Jelly aims to implement is RV32IC. This implements the base RISC-V integer instruction set, RV32I, as well as the “C” Compressed Instruction extension. In total, RV32I includes 47 instructions and the “C” extension includes an additional 16 instructions. I will therefore have to implement and verify a total of 63 instructions. This sounds like a lot, but do keep in mind the 16 compressed instructions have a 1-1 mapping with uncompressed instructions - so I only need to prove that they decode correctly.

The reason the “C” extension was added is that it's only some relatively simple extra logic in the decoder, and can significantly increase code density, leading to faster program execution.

We will also support Chapter 10, “Performance Counters”, specifically the CYCLE counter. This will be used for onboard performance timing, given we know the CPU’s f_{CLK} .

3.2 Pipeline

Unlike most processors, Jelly is *not* pipelined (although this is an extension task). As stated above, the goal for Jelly is to be as simple as possible, and although a simple 4-stage RISC pipeline is well documented, it may increase complexity and verification difficulty due to data hazards.

Instead, Jelly will perform the fetch, decode and execute operations on separate cycles. The completed instruction flow will look something like this:

Fetch (cycle 1)

Instruction memory is clocked, bringing the instruction at PC from the IMEM to chip registers. Although Jelly’s IMEM is backed by FPGA BRAM, it still takes at least one cycle to read the data in this BRAM.

Decode (cycle 2)

The instruction is decoded by combinatorial logic into its operands. This will probably involve decoding the instruction to ALU, LSU (load store unit) and BU (branch unit) operands. If the instruction failed to decode in this stage, the illegal instruction flag will be set, and Jelly will eventually jump to the fault handler (which exists at a fixed address in RAM). (TODO: maybe use handler jump table like AVR).

Execute (cycle 3)

In this stage, the ALU is run. Jelly will most likely also contain a barrel shifter, to make sure that we can execute any shift instruction on one clock cycle. This is mainly done for ease of implementation, but may also have performance benefits as well. If an instruction was not able to complete in one cycle, the execute stage will be re-run until the instruction has completed. If the instruction is a load-store instruction, then the load-store unit is asked to begin reading or writing to RAM. If the instruction specifically was a load instruction, then execute may need to be re-run to read the value from RAM. If the instruction was a branch, then the branch target is also computed by the ALU in this stage.

Writeback (cycle 4)

The results from the ALU are written back into the register file, and PC is set to the correct address based on flags and information set previously (e.g. jumps, branches, traps, etc).

3.3 Memory and privilege

The CPU will only implement the RISC-V Unprivileged ISA, as Jelly effectively shares its processing class with microcontrollers. To that end, there are currently no plans to add an MMU to the Jelly design.

Jelly implements a von Neumann architecture. Both the data memory and instruction memory will be backed by ECP5 block RAM (BRAM), which Lattice calls “sysMEM”. The exact size of this will depend on what sort of design resources the CPU takes, but I’m targeting at least 64 KiB of RAM and approximately 128 KiB of program memory. Theoretically we can target significantly more.

Although Jelly’s instruction memory is backed by BRAM during execution, the program is first copied from an external Winbond SPI flash module, to allow for simpler runtime programming without re-flashing the FPGA itself. This copy operation will be implemented as a separate CPU hardware peripheral that runs before boot, being completely invisible and inaccessible to the end-user. All it will do is use quad SPI fast-read mode to read the first 128 KiB from the flash chip using Winbond’s protocol, so should take only a few milliseconds after the PLL stabilises.

Note that the above SPI flash module will be separate and distinct SystemVerilog from the main CPU. This means that, should Jelly be synthesised for an ASIC in the future, this module could easily be removed or replaced with on-chip flash.

TODO: memory map

3.4 Interrupts

I don’t have massive plans in terms of the interrupt controller - just something that works enough to get UART in and out, and to handle system traps like an illegal instruction. This is somewhat of a departure from the

microcontroller design specification, however, it's worth noting that Jelly is not capable of being an actual MCU as it's missing other peripherals like GPIO. There are currently no plans for any sort of nested interrupt controller at all.

Jelly will most likely implement an interrupt system similar to the AVR. At the beginning of the program, an interrupt vector table will list jump addresses for specific types of interrupts, which the processor will jump to if an ISR or fault occurs.

3.5 Boot process

1. Power is sent to the FPGA, triggering its internal reset circuitry
2. FPGA loads Jelly's RTL from SPI flash or on-chip memory
3. Jelly RTL begins executing
4. Jelly initialises the OSCG oscillator and waits 1000 clocks for it to stabilise
5. SPI module copies first 128 KiB from Winbond SPI flash into BRAM
6. CPU state is reset
7. PC is set to **TODO: instruction memory addr** and the CPU begins executing instructions from there

3.6 Other microarchitecture details

The remaining details, such as what units the processor has and how it's organised on a lower level, will be decided while the thesis is in progress, as this will likely be an iterative process.

4 Tools and resources

4.1 HDL

Jelly will be implemented using the **SystemVerilog** hardware description language (HDL). This is because I already have familiarity with it, having started around November 2022, and it makes describing complex designs easier than plain Verilog or VHDL.

In addition, I have spent the last year developing "Slingshot", a SystemVerilog Language Server Protocol (LSP). This provides editor completion and diagnostics services, making SV a very productive language for me to edit in. Slingshot was specifically developed in anticipation of this thesis.

4.2 Simulation, verification and synthesis

For simulation and verification, I will depend exclusively on open source tools. These are perfectly fine for my use-case, and most importantly, free. The specific tools are:

- **Verilator**: An open-source, fast, cycle-accurate SystemVerilog simulator that works by transpiling SV to C++. It will be used extensively for verification. The only downside is it doesn't support verification methodologies like UVM, which is acceptable because I will be verifying the processor using a different method (see below).
- **Yosys**: An open-source EDA synthesis tool. Yosys is capable of synthesising SystemVerilog to a netlist, which will be used for FPGA implementation.
- **nextpnr**: An open-source place and route tool. Nextpnr has support for Lattice ECP5 FPGAs, and will be run after Yosys to place and route the synthesised netlist.
- **Catch2**: An open-source unit testing framework for C++. This will be used to create basic unit tests for the SystemVerilog code after Verilator has translated it to C++.

5 Verification plan

Ensuring that Jelly conforms to the RISC-V specification is arguably the most important task here. Jelly aims to be 100% compliant with the RISC-V RV32 specification, even in obscure edge cases, should they exist. To that end, I plan to verify Jelly in the following way, in order:

5.1 Verilator unit tests

Simple unit tests will be performed for each SystemVerilog module, for example the ALU. These will be written in C++ and use Verilator as the simulation platform. The unit tests themselves will use the Catch2 library. This is a good way to run an initial sanity check of the design, but is not sufficient for verifying an entire CPU. It's also useful during development to make sure that the device is functioning as intended.

5.2 RISC-V compliance test

In this test, the CPU will be run through the RISC-V Architecture Test SIG provided suite of test vectors. This will be a good initial test to prove that the design is functional, and should be very easy to use. These test suites are capable of auto-checking themselves, which will be very useful for detecting issues in the SystemVerilog code.

The RISC-V test suite that will be used is available from here:

<https://github.com/riscv-non-isa/riscv-arch-test>

5.3 Random instruction generation

In this process, a tool called **force-riscv** will be used to randomly generate combinations of valid RISC-V instructions. Each of these instructions will then be run through the Jelly CPU, which we aim to verify, and the official **Spike** RISC-V simulator, which we know is valid. We will compare the register state before and after each instruction, and make sure that Jelly is exactly identical to Spike.

This will be run repeatedly for a gruelling 12 hour period, and all inconsistencies against Spike will be treated as bugs and fixed.

Note that, based on some initial research, **force-riscv** appears to be a bit stubborn and difficult to use, so either may be skipped in favour of LibFuzzer (below) or using a custom, simpler RISC-V random instruction generator (although this may blow out the scope of the project too much).

6 Benchmarking

The main benchmark I intend to perform on Jelly is CoreMark and Dhrystone. For Dhrystone, I'll measure DMIPS, or Dhrystone million instructions per second, which is commonly used as a rough proxy for CPU performance.

Once synthesised for the FPGA, I will measure the total CoreMark score and CoreMark/MHz, and DMIPS, and compare against similar processors. In simulation, I will compare against an "ideal" 100 MHz clock speed, to gauge how the processor might perform on an ideal ASIC tapeout.

I also plan to run the RV32 port of FreeRTOS on the core and verify that context switching works as intended.

Unfortunately, I cannot guarantee a particular clock frequency as a deliverable, because this is extremely challenging to estimate at this early stage in the project. We can certainly target > 16 MHz as a reasonable goal, and we could theoretically reach up to 100 MHz with significant optimisation. The purpose of the benchmarking is to compare Jelly against existing industry designs, and to understand the processor's strengths and weaknesses.

7 Possible extension tasks

This section contains improvements that I will add to Jelly if time permits. They are ordered in the order that I would attempt them (simplest first).

Note: These are only suggestions for extension tasks. I may not attempt any of these at all, and definitely won't attempt all of them.

7.1 Implement “speculative pre-fetching”

Being a non-pipelined processor, Jelly fetches instructions on one cycle and executes them on the next. In this task, I'll implement a technique I've called “speculative pre-fetching” to potentially eliminate the fetch *and* decode cycle. In industry, this is usually called “instruction prefetch” or “cache prefetch”.

The instruction memory and instruction fetcher will be modified to read up to two instructions simultaneously (note this implies that the BRAM is at least dual-port, which the Lattice ECP5 does indeed support). The first instruction fetched is the actual instruction at \$PC, but the new *second* instruction fetched is whichever instruction the CPU believes it will most likely execute next. In the simple case of sequential code, this is just \$PC+4. I also plan to modify the instruction decoder to simultaneously decode two instructions at a time. This means that if the pre-fetcher is correct, two entire cycles can be skipped.

The skipping mechanism would run on the Writeback stage of the “pipeline”. When the pre-fetcher runs on the Fetch stage, it will leave a note indicating where it expects \$PC to point after Execute. On the Writeback stage, if the current position of \$PC matches the guess, then the CPU skips straight to Execute (since we already have Fetched and Decoded the next instruction). Note this means we cannot pre-fetch two instructions in a row. Since we jump straight to Execute, there aren't any cycles remaining to speculate about where the \$PC will change to. This is somewhat problematic, and we will need to benchmark whether or not this actually makes the CPU noticeably faster.

This approach will easily work in sequential code, but would always be wrong for branches. To fix this, we could implement a simple branch predictor (as is another extension task below). However, the branch prediction data structure may be complicated to implement, as each branch instruction would have to be “tagged” with an expected taken/not taken bit.

7.2 Add a fast “CLZ” instruction

In this task, we'll implement a part of the RISC-V Bit Manipulation instruction set, the Count Leading Zeroes (CLZ) instruction. This specific implementation will execute in a single cycle. CLZ is very useful for various types of algorithms and cryptography, and is a good example of how specialised hardware can be orders of magnitude faster than software.

7.3 Add “M” and “D” extensions

In this task, I'll add the RISC-V Multiply and Divide extensions. Of these, the “M” extension is the most important because multiply is usually more commonly used than divide.

It is currently undecided how this will be implemented and what type of performance to be expected from it. The goal is ideally to at least beat the software implementation of multiply. This may involve running the multiplier circuit at a faster clock speed than the main processor.

7.4 Implement a simple RISC pipeline

In this project, I would significantly overhaul Jelly from being a naive processor into a fully pipelined one, which opens the door for many further optimisations.

I plan to use the classic RISC pipeline, as it is the easiest to implement and the most well documented. The stages for this pipeline are: fetch, decode, execute, write-back. We could also look into implementing a shorter or longer pipeline.

This will involve a fair amount of work, because we will need to re-verify the entire processor to make sure that it handles hazards correctly. I plan to handle hazards simply by stalling, as this is an in-order machine.

As fun as this is, I would only really attempt this if I complete and verify the Jelly design significantly *ahead* of schedule (e.g. 6 months ahead of schedule), as it would be a lot of work.

7.5 Add a simple branch predictor

With the pipeline created, conditional branches may cause problems with stalling. Most modern processors use a *branch predictor*, a circuit that attempts to guess which way a branch will go.

The algorithm I plan to use is a 2-bit saturating counter (as documented on Wikipedia). This is one of the most simple branch predictors, but should actually show a performance increase.

7.6 SkyWater 130 nm VLSI tapeout using OpenLane

This is the most challenging subproject, as it is almost another paper in and of itself. It is very, very unlikely that this will actually be done. However, in this project, I'd tape out Jelly into an ASIC using open source tools. I would target SkyWater Technologies' 130 nm process node, using Google's free PDK. The synthesis tool would be a mix of Yosys for synthesis and OpenLane for place and route. This would be extremely complex to verify, as while open source tools *can* run a full tapeout, they don't have good power and thermal analysis tools. In addition, we would need to instantiate PDK-specific SRAM cells for the data memory of the CPU. Having this actually manufactured would be too expensive (and be somewhat of a waste for a processor as simple as Jelly), but the experience gained taping out the processor might be beneficial.